

Activity trace — accelerating matched SHD learning

ar073 · 19 July 2026 · Draft

Checkpoint CP-001

The next exploratory night is registered before any exp070 training result. This checkpoint has immutable private source SHA-256 prefix 5189910c09c9. Its checkpoint time is 2026-07-19 06:49:09.368 UTC, and it has no predecessor within this night. The raw snapshot is held read-only outside the repository. Hidden reasoning, tool payloads, injected environment material, credentials, private paths, addresses, and infrastructure identifiers are excluded. The visible attachment path is explicitly redacted; the request itself is reproduced verbatim.

Decision and action ledger

1. Pull request 52 is open and draft on the requested existing branch. Main is unchanged.
2. Exp069's complete epoch histories, firing diagnostics, gradients, rasters, split provenance, billing, and official-test prohibition were inspected.
3. A development-data-only audit rejected observation-window truncation as the likely bottleneck. The first registered candidate changes only the matched temporal resolution from 1 ms to 2 ms.
4. The ordered ladder, first-passing selection rule, promotion threshold, forty-epoch success criterion, kill criteria, one-seed policy, locked model comparison, and fresh paid-compute gate are fixed in ar072.
5. No exp070 runner, training result, cloud job, official-test access, or spend exists at this checkpoint. Pending work is the candidate-1 implementation, focused checks, local smoke, cost estimate, and paid-compute authority gate.

Privacy-sanitized visible transcript

User

2026-07-19 06:44:12.361 UTC

Role: user · Checkpoint: CP-001 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Files mentioned by the user:

/goal Design and execute the next exploratory SHD experiment on PR #52, aimed a...: (PRIVATE ATTACHMENT PATH REDACTED)

The attached pasted text file(s) contain the user's request. Read and act on that content.

My request for Codex:

Attached request

/goal Design and execute the next exploratory SHD experiment on PR #52, aimed at materially improving the matched COBA/PING baseline while preserving a fair scientific comparison.

Context:

- Continue from exp069 on branch night/spiking-heidelberg-digits/ar070.
- exp069 trained for 80 epochs and reached 43.87% held-out validation accuracy for COBA and 48.04% for PING.
- The additional epochs provided insufficient benefit. Use much shorter runs while developing, then return to a 40-epoch matched run after selecting a convincing configuration.
- One seed is sufficient for this exploratory phase. Do not introduce multiple seeds until there is a claim worth defending.
- Never access or tune against the official SHD test set during exploration.

Scientific objective:

Identify a mechanistically justified change that substantially improves SHD learning—ideally improving both cells while revealing whether PING offers an advantage over COBA.

Investigate likely bottlenecks in this order:

1. input temporal representation/resolution and shared input scaling;
2. initialization and optimization settings;
3. inhibitory recruitment, firing-rate balance, and gradient flow;
4. existing or carefully added neuron/synapse time-constant mechanisms;
5. only if justified by the diagnostics, a new matched readout structure.

Locked comparison:

- Preserve Dale constraints.
- Keep model capacity, hidden population sizes, E/I counts, connectivity, and recurrence structure fixed.
- COBA uses no voltage-gradient dampening.
- PING uses voltage-gradient dampening 1000.
- Keep dataset split, checkpoint-selection rule, batching, and all non-cell-specific settings matched.
- Any new mechanism, including a readout change, must be applied identically to COBA and PING unless the difference is intrinsic to their registered cell definitions.
- Do not change the scientific design after seeing results without recording the result and obtaining approval for the revised design.

Readout and tools/snn authority:

- Do not casually search over losses or readouts.
- Begin with exp069's existing readout and loss.
- A new readout structure is allowed only when the diagnostics provide a clear scientific justification and it is pre-registered before its results are observed.
- Keep the classification objective matched and avoid introducing a readout that obscures the recurrent-

cell comparison.

- You are authorized to make careful edits or additions to tools/snn when needed, including support for a justified new readout structure or neuron/synapse mechanism.
- Preserve backward compatibility, keep changes narrowly scoped, add focused tests, and document the engine commit anywhere it affects scientific interpretation.
- Stop and ask if the required engine change becomes broad, invasive, or changes the meaning of the COBA/PING comparison.

Protocol:

1. Inspect exp069's implementation, artifacts, learning curves, firing diagnostics, rasters, and epoch-by-epoch behavior. Diagnose the most likely bottleneck before changing anything.
2. Create and lock a timestamped mandate in the next arNNN article. Pre-register the hypothesis, allowed intervention sequence, selection rule, success criterion, and kill criterion.
3. Run a one-seed local smoke stage first.
4. Iterate using very short registered runs, preferably about 5 epochs and no more than 10 unless justified. Use only the held-out validation split.
5. Use a small sequential set of scientifically motivated interventions, normally changing one factor at a time. Do not conduct an unconstrained hyperparameter sweep.
6. Record every meaningful attempt, including null, failed, and killed attempts, before proceeding.
7. Promote a configuration only if its short-run curves, losses, gradients, and firing diagnostics provide credible evidence that it should outperform the exp069 trajectory.
8. Freeze the promoted configuration and run matched COBA and PING training for 40 epochs.
9. Compare the new run against exp069 at matched epochs and against exp069's final 80-epoch checkpoint. Do not claim improvement through selective checkpoint reporting.
10. Do not use the official SHD test set. Stop at the held-out-validation human-review gate.

Compute:

- Run focused local checks first.
- Estimate cloud runtime and cost before requesting paid-compute authorization.
- Do not assume an earlier RunPod budget applies to this experiment.
- If authorized, use at most one pod per cell, respect the approved ceiling, and reap every pod immediately.
- Do not run the full test suite on the 4 GB Hetzner host. Use focused lightweight tests and demolab build.

Deliverables:

- Implement the experiment through experiments/expNNN.py and narrowly scoped tools/snn changes where justified.
- Produce artifacts/data/expNNN with complete provenance, numbers.json, reproducer, learning curves, firing-rate and gradient diagnostics, and matched input/E/I rasters.
- Create a cold-readable writings/expNNN.typ whose numbers and figures are sourced from recorded artifacts.
- Maintain timestamped scientific records in the relevant arNNN article: decisions, hypotheses, actions, failures, anomalies, costs, commits, results, conclusions, and pending work.
- Continue the existing research activity-log article with meaningful timestamped, hash-linked checkpoints.

- Make focused commits for implementation, engine changes, and every meaningful scientific attempt, including killed attempts.
- Build and validate the Demolab publication.
- Push the existing branch and update PR #52 with useful milestone comments.
- Do not merge the PR or modify main.
- Finish at the human-review gate with a concise evidence summary, exact compute spend, validation performed, rendered article/report links, relevant commit links, and the PR link.

Assistant

2026-07-19 06:44:16.949 UTC

Role: assistant · Checkpoint: CP-001 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

I'll load the lab runbook and the attached `/goal`, then verify the current branch/PR and resume from the recorded exp069 state. I'll also use the GitHub publishing workflow when the experiment reaches a commit/push milestone.

Assistant

2026-07-19 06:44:20.058 UTC

Role: assistant · Checkpoint: CP-001 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The direct `demolab` command is not installed on this host, so I'm using the repository-prescribed `uvx demolab-cli docs` fallback before touching the experiment.

Assistant

2026-07-19 06:44:24.580 UTC

Role: assistant · Checkpoint: CP-001 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

This is an AUTORESEARCH continuation, so I'm loading that runbook and its programme rules, plus the attached mandate and GitHub commit/push instructions. I won't start paid compute until I've reconciled the registered design with the current branch and its existing records.

Assistant

2026-07-19 06:45:10.455 UTC

Role: assistant · Checkpoint: CP-001 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The programme is in "changes requested" state: you have explicitly approved a new exploratory night on the existing PR branch. I'm treating the attached goal as the mandate-gate approval, preserving its scientific terms verbatim, and keeping `main` untouched as requested; the branch commit will remain the pre-result anchor.

Assistant

2026-07-19 06:45:44.949 UTC

Role: assistant · Checkpoint: CP-001 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The exp069 evidence is intact locally, including full epoch histories and rasters. I'm now quantifying its early-epoch bottleneck, checking the next free IDs, and locating the existing private-trace custody pattern so the new night begins with a complete, sanitized checkpoint rather than a reconstructed summary.

Assistant

2026-07-19 06:47:51.951 UTC

Role: assistant · Checkpoint: CP-001 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The first input audit is testing two concrete failure modes on development data only: how much speech exp069 discarded after 1.0 seconds, and how much temporal information coarser bins would collapse. This precedes intervention selection and does not touch the official test split.

Assistant

2026-07-19 06:49:07.003 UTC

Role: assistant · Checkpoint: CP-001 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The diagnosis points to temporal discretization as the first clean intervention. Exp069 discarded only 0.017% of development events after 1 second, so extending duration is unlikely to move accuracy. A 2 ms bin collapses only 0.094% of events in a deterministic 1/8 development sample, while halving backpropagation depth; 4 ms already collapses 2.04% and is coarse relative to the 2 ms AMPA time constant. I'm preregistering 2 ms first, before observing any new training result.

Checkpoint CP-002

Candidate 1 is implemented and its local smoke gate passes. This checkpoint follows mandate anchor `cac3d07`. Its immutable private source has SHA-256 prefix `8b7c5612397e`, and its checkpoint time is `2026-07-19 07:04:04.503 UTC`.

Decision and action ledger

1. The exp070 runner changes only temporal resolution and short-run duration while reusing exp069's split, checkpoint selection, diagnostics, rasters, and cloud orchestration. `tools/snn` is unchanged.
2. A pre-smoke audit found that inherited backup/restore plumbing could copy an existing official-test cache even though it never evaluated it. Exp070 now redirects the engine to a development-only alias, closing that protocol ambiguity before training.
3. Focused Ruff, ty, compilation, and two-pod dry-run checks pass. The dry run created no pod.
4. Both 2 ms smoke cells are finite and active, use the exact exp069 split hashes, and report no skipped or non-finite updates. Spend remains 0 USD.
5. Pending work is a commit/push checkpoint and fresh authority for the matched five-epoch cloud pair. No paid compute may start before that approval.

Visible messages added in CP-002

Assistant

2026-07-19 06:50:08.850 UTC

Role: assistant · Checkpoint: CP-002 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

I've frozen the first private transcript checkpoint outside the repository as read-only, with SHA-256 prefix `5189910c09c9`. The public trace will redact the attachment's private path and preserve the message plus attached mandate verbatim; no credentials, tokens, addresses, environment values, or infrastructure identifiers will enter the article.

Assistant

2026-07-19 06:52:40.014 UTC

Role: assistant · Checkpoint: CP-002 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The mandate and CP-001 trace are written. I'm validating three things before the checkpoint commit: the trace JSON is syntactically sound, the publishable files contain no sensitive path/token/address patterns, and the entire Demolab collection renders with the new articles.

Assistant

2026-07-19 06:54:07.371 UTC

Role: assistant · Checkpoint: CP-002 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The build succeeds with 40 entries. It also refreshed every tracked PDF and emitted several stale untracked stub PDFs, which is normal build churn rather than experiment content. I'm narrowing the checkpoint to the two new article PDFs, the updated book, the articles, and the sanitized trace data; existing articles/PDFs remain untouched.

Assistant

2026-07-19 06:55:56.532 UTC

Role: assistant · Checkpoint: CP-002 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

CP-001 passes the full 40-entry build and the fail-closed sensitive-pattern scan. I'm committing only the new mandate, activity trace/data, their PDFs, and the updated book, then pushing the existing branch so PR #52's preview advances.

Assistant

2026-07-19 06:57:23.024 UTC

Role: assistant · Checkpoint: CP-002 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The first candidate needs no engine change: exp069 already exposes timestep, preserves the exact split, and records the required losses, gradients, rates, and checkpoints. I'm implementing exp070 as a narrow orchestration layer over that validated runner, with 2 ms hardcoded, a local smoke default, two short cloud cells only, and an automatic promotion/kill decision against the registered epoch-five baselines.

Assistant

2026-07-19 06:59:38.280 UTC

Role: assistant · Checkpoint: CP-002 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The focused checks exposed only two type-level literal assignments, now fixed. Before smoke, I caught a stricter protocol issue in the inherited plumbing: exp069 temporarily backed up any existing `(OFFICIAL

TEST CACHE PATH REDACTED)` while installing validation under the CLI's fixed filename. It never evaluated that file, but exp070's "no access" rule is stronger. I'm isolating the trainer's SHD alias in an experiment-specific directory so the official test file is neither copied nor opened.

Assistant

2026-07-19 07:00:27.614 UTC

Role: assistant · Checkpoint: CP-002 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The isolated runner now passes Ruff, `ty`, compilation, and a two-pod dry run; the dry run created nothing. I'm starting the registered local 128/128, two-epoch candidate-1 smoke. It will read only official training data to derive development train/validation subsets, and the engine will see those subsets in the exp070-specific alias directory.

Assistant

2026-07-19 07:01:42.047 UTC

Role: assistant · Checkpoint: CP-002 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Candidate 1's COBA smoke completed cleanly at 2 ms: finite two-epoch training, active 16 Hz excitatory firing, zero inhibitory activity as registered, and no update failures. PING is now running under dampening 1000 in the same isolated split.

Assistant

2026-07-19 07:04:01.063 UTC

Role: assistant · Checkpoint: CP-002 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The registered local smoke passed both cells. COBA finished active at 16.18 Hz E; PING at 4.27/15.00 Hz E/I; both had finite losses, exact exp069 split hashes, and no validation/test alias leakage. I'm promoting the smoke evidence into the committed artifact record, then checkpointing the runner and timestamped trace before asking for paid-compute authority.

Checkpoint CP-003

The cloud boundary now fails closed on scientific and billing integrity. This checkpoint follows smoke commit 135b82e. Its immutable private source has SHA-256 prefix 20a7b2d23008, and its checkpoint time is 2026-07-19 07:11:32.044 UTC.

Decision and action ledger

1. The expected first-candidate plus promoted-final path is approximately 2 USD. The requested hard ceiling is 10 USD across the entire registered ladder, with one pod per cell and two concurrent pods maximum. Authority has not yet been granted, so no dispatch occurred.
2. GitHub Actions reports the same ten repository-wide ty diagnostics in unchanged older experiments and tests. Exp070 produces no CI diagnostic, and those unrelated files remain out of scope.
3. Two focused tests pin the first-passing promotion gate: both cells must clear the accuracy and loss thresholds, and non-finite or saturated activity kills the candidate.

4. Artifact publication now refuses to proceed without finite authoritative provider spend and confirmation that zero pods remain. Ruff, focused `ty`, compilation, both tests, and the missing-ledger failure check pass.
5. Pending work remains fresh paid-compute authorization, the matched five-epoch pair, exact billing reconciliation, and the registered adaptive decision. Spend remains 0 USD.

Visible messages added in CP-003

Assistant

2026-07-19 07:07:24.169 UTC

Role: assistant · Checkpoint: CP-003 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The local checkpoint is clean: focused Ruff, ``ty``, compilation, JSON validation, sensitive-pattern scan, exact split/config assertions, and the live 40-entry Demolab rebuild all pass. I'm committing this smoke attempt now; ``tools/snn`` remains unchanged and RunPod spend is still exactly \$0.

Assistant

2026-07-19 07:08:15.643 UTC

Role: assistant · Checkpoint: CP-003 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Candidate 1 is ready for cloud testing.

Estimated RunPod cost:

- Five-epoch 2 ms pair: about \$0.40 expected
- If promoted, forty-epoch 2 ms pair: about \$1.50 expected
- Expected combined path: about \$2
- Proposed hard ceiling across all registered candidates and final run: \$10
- Maximum concurrency: two pods, exactly one per cell
- Backstops: one hour for short runs, two hours for final runs
- All pods reaped immediately

Current spend is \$0. Smoke checkpoint: ``135b82e``; PR #52 is updated.

Please authorize RunPod spending up to the \$10 hard ceiling for exp070.

Assistant

2026-07-19 07:08:32.804 UTC

Role: assistant · Checkpoint: CP-003 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Paid compute remains gated pending your explicit authorization, so I won't dispatch anything. I'll use this continuation to harden exp070's remaining registered paths, verify the pushed preview/checks, and ensure cost reconciliation and the eventual 40-epoch promotion cannot proceed with missing evidence.

Assistant

2026-07-19 07:08:46.471 UTC

Role: assistant · Checkpoint: CP-003 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

PR #52 is still draft and unmerged; its preview is updating. The new lint check is red, so I'm using the GitHub CI-diagnosis workflow to verify whether exp070 introduced anything or whether this is the already-documented repository-wide type debt. I'll only change files in scope if the logs point to exp070.

Assistant

2026-07-19 07:09:26.515 UTC

Role: assistant · Checkpoint: CP-003 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

CI confirms exp070 adds no diagnostic: the lint failure is the same ten pre-existing `ty` errors in unchanged exp041/048/054/065 and their older tests. I'm not expanding scope to fix them. One in-scope hardening remains before dispatch: make short-result publication fail closed unless exact provider billing and zero active pods are recorded, and enforce the registered validation-rate saturation checks explicitly.

Assistant

2026-07-19 07:11:19.187 UTC

Role: assistant · Checkpoint: CP-003 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The hardening checks now pass: two focused promotion-gate tests verify that both cells must clear accuracy/loss criteria and that non-finite or saturated activity kills the candidate. Publication also fails immediately if the exact billing ledger or zero-pod confirmation is absent. I'm recording this CI/provenance checkpoint while the paid-compute gate remains closed.

Checkpoint CP-004

Fresh paid-compute authority is explicit and bounded. This checkpoint follows hardening commit 2bc22d3. Its immutable private source has SHA-256 prefix c20363a95684, and its checkpoint time is 2026-07-19 07:21:26.366 UTC.

Decision and action ledger

1. The scientist authorized a 10 USD hard total ceiling across exp070, with at most two concurrent pods and exactly one cell per pod.
2. Short candidates carry one-hour pod backstops. A promoted forty-epoch pair may use two-hour backstops. Every pod must be reaped immediately.
3. Before this checkpoint the fleet was empty, spend was 0 USD, the branch was clean at 2bc22d3, and the pull request remained draft and unmerged.
4. The next action is limited to the registered five-epoch 2 ms pair. No later candidate or final pair is dispatched before applying the registered result-dependent decision.

Visible messages added in CP-004

Assistant

2026-07-19 07:14:17.507 UTC

Role: assistant · Checkpoint: CP-004 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Exp070 is fully hardened and ready to dispatch from `2bc22d3`; current spend is \$0.

Requested authorization: RunPod up to a \$10 hard total ceiling, maximum two concurrent pods (one COBA, one PING), with one-hour short-run and two-hour final-run backstops. Expected spend is about \$2 if candidate 1 is promoted.

Please explicitly authorize this paid compute.

Assistant

2026-07-19 07:14:28.941 UTC

Role: assistant · Checkpoint: CP-004 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The only remaining gate is unchanged: paid RunPod authority has not been granted, so dispatching would violate the registered mandate. I'm doing a final safety check for zero active pods, clean remote state, and the latest PR preview/check status; if those are intact, I must formally stop at this authority blocker rather than spend without permission.

Assistant

2026-07-19 07:14:44.983 UTC

Role: assistant · Checkpoint: CP-004 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Blocked only on fresh paid-compute authorization.

Safety state:

- Active pods: 0
- Spend: \$0
- Branch clean and pushed at `2bc22d3`
- PR #52 remains draft and unmerged
- Candidate-1 smoke and focused checks pass
- CI lint failure remains unrelated pre-existing type lint

Authorize RunPod up to the \$10 hard ceiling to resume.

User

2026-07-19 07:21:02.887 UTC

Role: user · Checkpoint: CP-004 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

authorized

Assistant

2026-07-19 07:21:09.728 UTC

Role: assistant · Checkpoint: CP-004 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Authorization received for the proposed exp070 envelope: \$10 hard total ceiling, at most two concurrent pods with one per cell, one-hour short-run backstops and two-hour final-run backstops, with immediate

reaping. I'll checkpoint this authority first, verify the fleet is empty, then dispatch only the registered five-epoch 2 ms COBA/PING pair from `2bc22d3`.

Checkpoint CP-005

Only the registered candidate-1 pair is dispatched. This checkpoint follows authority commit `cd82023`. Its immutable private source has SHA-256 prefix `3b4b22492a5c`, and its checkpoint time is `2026-07-19 07:23:43.025 UTC`.

Decision and action ledger

1. The dry run found zero existing pods and exactly two jobs: one COBA and one PING cell at 2 ms for five epochs.
2. Two 5090 pods were created at 0.99 USD/hour each, pinned to `cd82023`, with one-hour self-removal backstops. The post-dispatch fleet count was exactly two, so worst-case exposure for this attempt is 1.98 USD.
3. Pod identifiers and other infrastructure details remain only in the private trace. The public record retains names, counts, rates, commit, and bounds.
4. Pending work is self-termination, collection, exact provider billing, artifact publication, and the preregistered promotion decision.

Visible messages added in CP-005

Assistant

2026-07-19 07:22:57.584 UTC

Role: assistant · Checkpoint: CP-005 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The authority checkpoint is pushed at `cd82023`; the scientific runner is unchanged from the tested `2bc22d3` state. I'm performing the final dry run and fleet check now. If it still resolves to exactly two jobs and zero existing pods, I'll dispatch the five-epoch pair under the one-hour backstops.

Assistant

2026-07-19 07:23:13.646 UTC

Role: assistant · Checkpoint: CP-005 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The dry run is exact: zero existing pods, then one COBA and one PING job on two 5090 pods, with no extra work. Worst-case exposure for this short pair is \$1.98, below the \$10 cumulative ceiling. Dispatching now from `cd82023`; I'll verify the fleet count immediately afterward.

Assistant

2026-07-19 07:23:36.032 UTC

Role: assistant · Checkpoint: CP-005 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The registered pair is live: exactly two pods at \$0.99/hour each, one named COBA and one PING, pinned to `cd82023`; no third pod exists. The one-hour guards cap this attempt at \$1.98. I'm checkpointing the dispatch without publishing pod identifiers, then monitoring for self-termination before collection and billing reconciliation.

Checkpoint CP-006

Candidate 1 is killed after the registered five-epoch comparison. This checkpoint follows dispatch commit `1bba764`. Its immutable private source has SHA-256 prefix `cf6b7ebf703b`, and its checkpoint time is `2026-07-19 07:34:41.277 UTC`.

Decision and action ledger

1. Both pods self-terminated about eight minutes after dispatch; the active fleet is zero and the complete cell outputs were collected.
 2. COBA reached 23.53% validation accuracy at epoch five, 4.41 percentage points below `exp069`, while cross-entropy worsened by 0.206 to 2.678.
 3. PING reached 27.70%, 4.90 points below `exp069`, while cross-entropy worsened by 0.148 to 2.592.
 4. Both cells were finite and active with no skipped or non-finite updates, so this is a scientific kill rather than a plumbing failure. The 2 ms change fails every registered promotion requirement.
 5. Exact provider billing has not yet appeared in the billing history. The publication gate remains closed until the authoritative spend is recorded; no estimate is substituted.
- Candidate 2 has not been implemented or run.

Visible messages added in CP-006

Assistant

2026-07-19 07:25:21.662 UTC

Role: assistant · Checkpoint: CP-006 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Dispatch checkpoint ``1bba764`` is pushed; I also corrected the PR comment's literal commit placeholder without editing history. Both jobs remain active. I'm switching to a single bounded monitor that reports only timestamp, job names, and fleet count; it will exit as soon as both pods self-terminate.

Assistant

2026-07-19 07:26:21.152 UTC

Role: assistant · Checkpoint: CP-006 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Both candidate jobs remain active at 07:25:59 UTC. Fleet count is still exactly two, estimated exposure is under \$0.10, and the one-hour safeguards remain intact. No duplicate process or additional pod has been created.

Assistant

2026-07-19 07:28:07.372 UTC

Role: assistant · Checkpoint: CP-006 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The pair is still healthy and bounded at 07:27:32 UTC. This is consistent with cold image startup plus staging before the five short epochs; estimated combined exposure is about \$0.15. I'm continuing the same monitor.

Assistant

2026-07-19 07:30:33.592 UTC

Role: assistant · Checkpoint: CP-006 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Both pods remain present at 07:30:07 UTC, about seven minutes after dispatch. Estimated combined exposure is about \$0.23. That remains within the expected startup-plus-short-run window and far below the \$1.98 attempt cap.

Assistant

2026-07-19 07:31:15.947 UTC

Role: assistant · Checkpoint: CP-006 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Both training pods self-terminated by 07:31:10 UTC, about eight minutes after dispatch; active fleet is now zero. I'm collecting the shared-volume outputs with the bounded collector, then I'll verify both five-epoch records and reconcile exact provider billing before allowing publication or any next candidate.

Assistant

2026-07-19 07:33:04.359 UTC

Role: assistant · Checkpoint: CP-006 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The first registered cloud candidate has finished and is decisively killed: moving to 2 ms reduced held-out accuracy and worsened cross-entropy in both COBA and PING. I'm reconciling exact provider billing now, then I'll publish and checkpoint this failed attempt before touching candidate 2.

Assistant

2026-07-19 07:34:22.996 UTC

Role: assistant · Checkpoint: CP-006 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Provider billing is lagging behind pod termination, so I'm not inventing an "exact" cost. The fleet is confirmed empty; while the authoritative entries settle, I'm preserving the candidate's outputs and preparing the timestamped killed-attempt record.

Checkpoint CP-007

Candidate 2 is implemented and passes its local smoke gate. This checkpoint follows the killed-attempt commit **9e1385a**. Its immutable private source has SHA-256 prefix **b12370990974**, and its checkpoint time is **2026-07-19 07:47:15.296 UTC**.

Decision and action ledger

1. Candidate 2 restores 1 ms and changes only the matched shared input-weight mean from 0.9 to 1.2. Candidate 1's outputs remain separately addressable.
2. Focused Ruff, type, compile, and promotion-gate checks pass. `tools/snn` is unchanged.
3. The required 128-training/128-validation two-epoch smoke passes in both cells. COBA ends at 26.79 Hz E; PING ends at 3.24 Hz E and 17.00 Hz I. Losses and gradients are finite and no update is skipped.
4. The exact registered split hashes are retained and the development-only alias contains no official-test file. The two-job cloud dry run creates no pod and resolves to exactly one job per cell.

5. Candidate-1 provider billing is still delayed. The fleet is empty and candidate 2 has not been dispatched at this checkpoint.

Visible messages added in CP-007

Assistant

2026-07-19 07:37:37.732 UTC

Role: assistant · Checkpoint: CP-007 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The killed candidate is now committed and pushed at `9e1385a`, with its timestamped activity checkpoint and PR milestone comment. No later candidate has started. I'm waiting only for the provider's authoritative billing rows so this attempt can be published with exact spend, then I'll advance to the preregistered 1.2 input-scale smoke.

Assistant

2026-07-19 07:41:19.944 UTC

Role: assistant · Checkpoint: CP-007 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Candidate 2 is now implemented exactly as registered: 1 ms is restored and only the shared input-weight mean changes from 0.9 to 1.2. The runner can still reproduce and publish the killed 2 ms attempt without mixing its files. Focused lint, type, compile, and promotion-gate tests pass; I'm running the required local 128/128 two-epoch smoke now.

Assistant

2026-07-19 07:46:43.509 UTC

Role: assistant · Checkpoint: CP-007 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The candidate-2 smoke passes: both cells are finite and active, use the exact registered 7,340/816 split hashes, and the development-only alias contains no official-test file. COBA ends at 26.79 Hz E; PING at 3.24 Hz E and 17.00 Hz I, with no skipped/non-finite updates. I'm checkpointing this smoke and verifying the two-job dry run before any paid dispatch.

Checkpoint CP-008

Only the registered candidate-2 pair is dispatched. This checkpoint follows smoke commit `04d5a28`. Its immutable private source has SHA-256 prefix `a0717f4a44f2`, and its checkpoint time is 2026-07-19 07:50:04.550 UTC.

Decision and action ledger

1. The clean dry run found zero existing pods and exactly two candidate-2 jobs: one COBA and one PING cell at 1 ms with shared input mean 1.2.
2. Two 5090 pods were created at 0.99 USD/hour each, pinned to `04d5a28`, with one-hour self-removal backstops. The verified post-dispatch fleet count is exactly two.
3. Candidate 1's exact billing is still delayed, but its hard bound is 1.98 USD. The two short attempts therefore have a cumulative hard bound of 3.96 USD, below the authorized 10 USD ceiling.
4. Infrastructure identifiers remain only in the immutable private trace. Pending work is self-termination, collection, provider reconciliation, and the locked promotion decision.

Visible messages added in CP-008

Assistant

2026-07-19 07:49:25.308 UTC

Role: assistant · Checkpoint: CP-008 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Candidate 2 is committed and pushed at `04d5a28`. Exact candidate-1 billing is still absent from the provider history, but its one-hour guard makes the hard cumulative exposure at most \$1.98; adding this registered short pair can raise that bound only to \$3.96, safely below the authorized \$10 ceiling. I'm posting the smoke milestone, rerunning the clean dry run, then dispatching exactly one pod per cell.

Assistant

2026-07-19 07:49:57.899 UTC

Role: assistant · Checkpoint: CP-008 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The candidate-2 pair is live from clean commit `04d5a28`: exactly two 5090 pods at \$0.99/hour, one COBA and one PING, with no third pod. The one-hour safeguards put the cumulative worst-case bound at \$3.96 across both short attempts. I'm checkpointing this dispatch without publishing infrastructure identifiers, then monitoring until the fleet returns to zero.

Checkpoint CP-009

Candidate 2 is killed after the locked epoch-five comparison. This checkpoint follows dispatch commit 90575b9. Its immutable private source has SHA-256 prefix c55f4048b3c7, and its checkpoint time is 2026-07-19 08:06:39.010 UTC.

Decision and action ledger

1. PING self-terminated first and COBA followed by 08:05:03 UTC. The fleet is zero and both complete records were collected.
2. COBA reaches 27.94% at epoch five, exactly matching exp069, while cross-entropy is 0.0168 worse. PING reaches 33.70%, only 1.10 percentage points above exp069, while cross-entropy is 0.00153 worse.
3. Both cells are finite and active with no skipped or non-finite updates. Candidate 2 therefore fails the efficacy gate cleanly; it is not promoted.
4. Exact provider rows for both completed attempts remain delayed. No estimate is published as exact. Candidate 3 is the sole remaining registered short attempt and has not yet been implemented.

Visible messages added in CP-009

Assistant

2026-07-19 07:53:03.411 UTC

Role: assistant · Checkpoint: CP-009 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Candidate 2 remains exactly two pods at 07:52:40 UTC, roughly three minutes after dispatch. The jobs are still within normal cold-start/training time, with estimated combined exposure under \$0.12 and the one-hour guards intact.

Assistant

2026-07-19 07:54:48.231 UTC

Role: assistant · Checkpoint: CP-009 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Both jobs remain present at 07:54:13 UTC, about four and a half minutes after dispatch. Fleet count is still exactly two, estimated combined exposure is about \$0.15, and no duplicate work has appeared.

Assistant

2026-07-19 07:56:34.593 UTC

Role: assistant · Checkpoint: CP-009 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

At 07:56:17 UTC the pair is still bounded and healthy from the fleet view, about seven minutes after dispatch. Estimated combined exposure is about \$0.23; both one-hour safeguards remain active.

Assistant

2026-07-19 07:58:55.750 UTC

Role: assistant · Checkpoint: CP-009 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The candidate remains at exactly two pods at 07:58:21 UTC, about nine minutes after dispatch. This 1 ms run is expected to take longer than the killed 2 ms attempt; exposure is still only about \$0.30 combined.

Assistant

2026-07-19 08:00:36.922 UTC

Role: assistant · Checkpoint: CP-009 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Both jobs remain active at 08:00:25 UTC, about eleven minutes after dispatch, with no fleet anomaly. Combined exposure is about \$0.36—still far below both the attempt backstop and total authorization.

Assistant

2026-07-19 08:02:25.380 UTC

Role: assistant · Checkpoint: CP-009 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

PING has self-terminated successfully by 08:01:57 UTC; COBA is the sole remaining pod. Fleet count is one, so spend is now accruing at only \$0.99/hour while COBA finishes under its guard.

Assistant

2026-07-19 08:04:10.524 UTC

Role: assistant · Checkpoint: CP-009 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

COBA remains the only active job at 08:04:01 UTC, roughly fifteen minutes after dispatch. This is still within the expected 1 ms five-epoch window; total candidate-2 exposure is approximately \$0.48 so far.

Assistant

2026-07-19 08:05:11.543 UTC

Role: assistant · Checkpoint: CP-009 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

COBA self-terminated by 08:05:03 UTC; the active fleet is now zero. Candidate 2 took about fifteen and a half minutes wall-clock, with estimated spend around \$0.49 pending authoritative billing. I'm collecting the outputs now and will apply the locked epoch-five promotion gate before any further run.

Assistant

2026-07-19 08:06:15.306 UTC

Role: assistant · Checkpoint: CP-009 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Candidate 2 is also killed by the preregistered gate. At epoch five, COBA exactly matches exp069 accuracy (27.94%, +0.00 pp) while CE is 0.0168 worse; PING improves only 1.10 pp to 33.70% while CE is 0.00153 worse. Both cells are finite and active with zero skipped/non-finite updates, so this is another clean scientific failure. Candidate 3—the isolated learning-rate change—is now the only remaining registered short attempt.

Checkpoint CP-010

The final registered short candidate passes its local smoke gate. This checkpoint follows the candidate-2 kill commit `3ac37b4`. Its immutable private source has SHA-256 prefix `1035934cfc77`, and its checkpoint time is `2026-07-19 08:14:25.022 UTC`.

Decision and action ledger

1. Candidate 3 restores the baseline 1 ms temporal resolution and shared input mean 0.9, changing only the matched Adam learning rate to 0.001.
2. Focused Ruff, type, compile, and promotion-gate tests pass. `tools/snn` is unchanged and no candidates are combined.
3. Both 128/128 two-epoch smoke cells pass with finite activity and no skipped updates. COBA ends at 26.13 Hz E; PING at 3.36 Hz E and 19.00 Hz I.
4. Exact split hashes are retained, the staged alias has no official-test file, and the cloud dry run creates no pod while resolving to one job per cell. Paid dispatch remains pending this checkpoint.

Visible messages added in CP-010

Assistant

2026-07-19 08:08:26.054 UTC

Role: assistant · Checkpoint: CP-010 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The second killed attempt is committed and pushed at ``3ac37b4``, with its timestamped activity record. I'm now implementing the final registered short candidate: restore input mean 0.9, retain 1 ms, and change only the matched Adam learning rate from 0.0004 to 0.001. No combinations or post-result tuning are being introduced.

Assistant

2026-07-19 08:10:32.854 UTC

Role: assistant · Checkpoint: CP-010 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Candidate 3's COBA smoke is progressing normally with the isolated 0.001 learning rate: finite loss, active E population, and no skipped updates through the second epoch. I'm letting the same bounded local smoke finish both cells.

Assistant

2026-07-19 08:14:10.988 UTC

Role: assistant · Checkpoint: CP-010 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Candidate 3 passes its local smoke gate: COBA ends at 26.13 Hz E; PING at 3.36 Hz E and 19.00 Hz I, with finite training, exact split hashes, and no official-test alias file. The clean dry run resolves to exactly two jobs and creates none. I'm checkpointing this final short candidate before dispatch.

Checkpoint CP-011

Only the final registered short pair is dispatched. This checkpoint follows smoke commit `e9d0332`. Its immutable private source has SHA-256 prefix `c49ca5689e35`, and its checkpoint time is 2026-07-19 08:16:37.476 UTC.

Decision and action ledger

1. The clean dry run found zero pods and exactly two candidate-3 jobs. The live dispatch created one COBA and one PING 5090 pod at 0.99 USD/hour each.
2. Both jobs are pinned to `e9d0332` and carry one-hour self-removal backstops. The verified fleet count is exactly two.
3. Using unreconciled one-hour maxima for all three short pairs gives a 5.94 USD cumulative hard bound, still below the authorized 10 USD ceiling.
4. Infrastructure identifiers remain private. Pending work is termination, collection, billing reconciliation, and the locked promotion decision.

Visible messages added in CP-011

Assistant

2026-07-19 08:16:08.625 UTC

Role: assistant · Checkpoint: CP-011 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Candidate 3 is committed and pushed at ``e9d0332``. Even using the unreconciled one-hour maxima for both completed attempts, this final short pair can raise cumulative exposure only to \$5.94, below the \$10 ceiling. I'm posting the smoke milestone and dispatching exactly the registered two-job pair from the clean commit.

Assistant

2026-07-19 08:16:32.029 UTC

Role: assistant · Checkpoint: CP-011 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The final short candidate is live from ``e9d0332``: exactly two 5090 pods at \$0.99/hour, one per cell, and no third pod. One-hour backstops hold the cumulative worst-case exposure at \$5.94. I'm recording the dispatch checkpoint, then monitoring to zero before collection and the final ladder decision.

Checkpoint CP-012

All registered short candidates are killed and the ladder terminates. This checkpoint follows dispatch commit **768af93**. Its immutable private source has SHA-256 prefix **a58a4f6c4400**, and its checkpoint time is **2026-07-19 08:38:08.255 UTC**.

Decision and action ledger

1. Both candidate-3 pods self-terminated, the fleet returned to zero, and the complete outputs were collected before evaluation.
2. At epoch five, COBA is 1.84 points below exp069 with cross-entropy 0.0137 worse. PING is 1.59 points above exp069 with cross-entropy 0.1148 better, but misses the locked +3-point threshold.
3. Both cells are finite and active with no skipped or non-finite updates. Candidate 3 is killed; all three candidates are now killed and the protocol forbids a forty-epoch run.
4. Candidate 1's authoritative provider charge is 0.251047 USD. Later billing rows remain unsettled, so exact cumulative spend is still pending. The candidate-1 artifacts and cold-readable experiment report are published in draft form while reconciliation continues.

Visible messages added in CP-012

Assistant

2026-07-19 08:19:53.646 UTC

Role: assistant · Checkpoint: CP-012 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Candidate 3 remains exactly two pods at 08:19:37 UTC, about three minutes after dispatch. No duplicate work is present; estimated combined exposure is under \$0.12 and the one-hour guards remain intact.

Assistant

2026-07-19 08:22:29.546 UTC

Role: assistant · Checkpoint: CP-012 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Both final short jobs remain active at 08:21:41 UTC, about five minutes after dispatch. The fleet is unchanged at two and combined exposure is approximately \$0.18.

Assistant

2026-07-19 08:24:28.746 UTC

Role: assistant · Checkpoint: CP-012 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

At 08:23:44 UTC the pair is still exactly two pods, about seven and a half minutes after dispatch. This remains within the expected 1 ms runtime; estimated combined exposure is about \$0.25.

Assistant

2026-07-19 08:26:15.724 UTC

Role: assistant · Checkpoint: CP-012 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Both jobs remain active at 08:25:48 UTC, about nine and a half minutes after dispatch, with no fleet anomaly. Combined exposure is roughly \$0.31.

Assistant

2026-07-19 08:28:00.827 UTC

Role: assistant · Checkpoint: CP-012 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

PING self-terminated by 08:27:52 UTC; COBA is the only remaining pod. The fleet has dropped to one normally, so only \$0.99/hour is accruing while COBA completes.

Assistant

2026-07-19 08:30:28.548 UTC

Role: assistant · Checkpoint: CP-012 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

COBA remains the sole active job at 08:29:55 UTC, about fourteen minutes after dispatch. Total candidate-3 exposure is approximately \$0.46 so far, still comfortably bounded.

Assistant

2026-07-19 08:31:34.627 UTC

Role: assistant · Checkpoint: CP-012 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

COBA self-terminated by 08:31:28 UTC and the fleet is zero. The final short pair took about fifteen minutes wall-clock, with estimated spend around \$0.50 pending billing. I'm collecting now; if either cell misses the locked gate, the ladder ends with no forty-epoch run.

Assistant

2026-07-19 08:32:28.480 UTC

Role: assistant · Checkpoint: CP-012 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Candidate 3 is killed, so the preregistered ladder ends and no forty-epoch run is allowed. COBA falls 1.84 pp to 26.10% with CE 0.0137 worse. PING improves 1.59 pp to 34.19% and CE improves by 0.1148, but still misses the required +3 pp; both cells had to pass. All three candidates therefore fail the shared promotion rule cleanly.

Assistant

2026-07-19 08:34:04.678 UTC

Role: assistant · Checkpoint: CP-012 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The ladder's scientific outcome is now fixed: all three candidates are killed, so the forty-epoch stage is prohibited. Candidate 1's authoritative billing has arrived at exactly \$0.251047; candidates 2 and 3 are still pending provider settlement. I've published candidate 1's full raw outputs and am extending the artifact layout so every killed attempt retains its own curves, activity diagnostics, and matched rasters in the final report.

Checkpoint CP-013

All work independent of final provider settlement is complete. This checkpoint follows aggregation commit 971e253. Its immutable private source has SHA-256 prefix bd61a8aa2b24, and its checkpoint time is 2026-07-19 09:14:12.307 UTC.

Decision and action ledger

1. Candidate 2's authoritative provider charge settled at 0.13992502 USD, bringing the exact two-attempt subtotal to 0.39097231 USD. Its complete raw data, curves, activity panels, and rasters are published.
2. Candidate 3's curves, activity diagnostics, and matched rasters were rendered directly from its already collected records while billing remains pending. All three rasters passed visual inspection.
3. Cross-attempt validation confirms five finite epochs per cell, zero skipped or non-finite updates, byte-identical paired inputs, no official-test file, and zero active pods.
4. The 41-entry Demolab build passes without stubs. GitHub preview and tests pass; lint remains red only from previously documented unrelated type debt.
5. The only pending work is candidate-3 authoritative billing, final atomic publication, cumulative cost insertion, rebuild, push, and review gate.

Visible messages added in CP-013

Assistant

2026-07-19 08:43:32.389 UTC

Role: assistant · Checkpoint: CP-013 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Artifact integrity checks pass across all three attempts: five complete finite epochs per cell, zero skipped/non-finite updates, and byte-identical matched inputs within each COBA/PING pair. The 1 ms candidates also share the same input hash. No official-test file appears in the staged or published experiment trees, and the fleet remains empty.

Assistant

2026-07-19 08:46:07.354 UTC

Role: assistant · Checkpoint: CP-013 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The provider has not posted either later pair's billing rows yet; this matches the roughly one-hour settlement lag seen for candidate 1. I'm keeping the report explicitly marked pending rather than substituting runtime estimates, while the zero-pod state remains verified.

Assistant

2026-07-19 08:48:21.519 UTC

Role: assistant · Checkpoint: CP-013 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Billing is still unsettled at 08:48 UTC. Nothing scientific or computational is running—the only remaining work is provider reconciliation, final per-attempt artifact publication, the complete report build, and review-gate handoff.

Assistant

2026-07-19 08:50:35.831 UTC

Role: assistant · Checkpoint: CP-013 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

No provider rows yet at 08:50 UTC; candidate 1's observed settlement pattern suggests candidate 2 should appear shortly after 09:05 UTC and candidate 3 around 09:31 UTC. The hard \$10 ceiling is

already safe—actual runtime exposure for all three pairs is under roughly \$1.25—but I'll report only authoritative amounts as exact.

Assistant

2026-07-19 08:51:57.434 UTC

Role: assistant · Checkpoint: CP-013 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Candidate 2's authoritative rows have now settled: COBA \$0.07004597 and PING \$0.06987905, totaling exactly \$0.13992502. I'm publishing its complete raw data and candidate-specific curves/rasters now; candidate 3's rows remain pending.

Assistant

2026-07-19 08:55:07.587 UTC

Role: assistant · Checkpoint: CP-013 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Candidate 2 is fully published and the live preview URLs return HTTP 200. Candidate 3 billing is the sole remaining external dependency; its scientific data and zero-pod state are already complete.

Assistant

2026-07-19 08:57:26.559 UTC

Role: assistant · Checkpoint: CP-013 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

No candidate-3 provider rows as of 08:57 UTC. The exact reconciled subtotal for candidates 1 and 2 is \$0.39097231; the final total remains deliberately unstated until both last-pair rows arrive.

User

2026-07-19 08:58:27.627 UTC

Role: user · Checkpoint: CP-013 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

status and eta?

Assistant

2026-07-19 08:58:37.850 UTC

Role: assistant · Checkpoint: CP-013 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Science and compute are complete: all three candidates were cleanly killed, no 40-epoch run is permitted, all pods are reaped, and candidates 1–2 are fully published.

Remaining: candidate 3's delayed provider billing rows, then its artifact publication, final report/activity checkpoint, build, push, and preview verification. ETA is about 25–40 minutes: likely 15–25 minutes for billing settlement, then roughly 10–15 minutes to finalize and validate. Exact reconciled spend so far is \$0.39097231; actual total will remain well below \$1.00.

Assistant

2026-07-19 09:00:51.269 UTC

Role: assistant · Checkpoint: CP-013 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

At 09:00 UTC the final billing rows are still pending; there is no running compute. The report, candidate decisions, and preview are already in place, so settlement remains the only item controlling the ETA.

Assistant

2026-07-19 09:03:03.542 UTC

Role: assistant · Checkpoint: CP-013 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Still waiting on provider settlement at 09:03 UTC. The previous pair's rows arrived about 46 minutes after its fleet reached zero; candidate 3 is now around 31 minutes past zero, so the current 15–25 minute billing estimate remains reasonable.

Assistant

2026-07-19 09:06:22.857 UTC

Role: assistant · Checkpoint: CP-013 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

No change at 09:06 UTC: zero pods, final scientific verdict fixed, candidate-3 billing rows pending. Once they appear, the remaining local publication/build/push sequence is about 10–15 minutes.

Assistant

2026-07-19 09:09:36.494 UTC

Role: assistant · Checkpoint: CP-013 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Provider settlement is still the only outstanding item at 09:09 UTC. There has been no additional spend since 08:31:28 UTC, when the final pod self-terminated.

User

2026-07-19 09:11:19.371 UTC

Role: user · Checkpoint: CP-013 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

whats the hangup?

Assistant

2026-07-19 09:11:29.817 UTC

Role: assistant · Checkpoint: CP-013 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The experiment itself is done. The hangup is only RunPod's delayed billing API: the final two pod charges haven't appeared yet, and you asked for exact spend, so I'm refusing to substitute an estimate.

Current state:

- All three candidates are killed; no 40-epoch run

- All outputs collected
- Active pods: zero since 08:31 UTC
- Exact settled subtotal: \$0.39097231
- Final pair's actual runtime exposure: roughly \$0.50, but not yet authoritative
- After its billing rows arrive: about 10–15 minutes to publish, build, push, and hand off

I'll check whether the rows landed during the interruption now.

Assistant

2026-07-19 09:11:38.842 UTC

Role: assistant · Checkpoint: CP-013 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Confirmed: fleet is still empty, and both final billing rows are still absent. This is provider-side accounting lag, not a training, artifact, or code failure. I'll poll less noisily and resume immediately when the authoritative charges appear.

User

2026-07-19 09:11:58.089 UTC

Role: user · Checkpoint: CP-013 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

do the other stuff in the meantime

Assistant

2026-07-19 09:12:20.336 UTC

Role: assistant · Checkpoint: CP-013 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Agreed. I'll stop treating billing as the critical path and finish everything independent of it now: render candidate 3's curves/rasters from collected data, complete the three-attempt report layout, run artifact and privacy validation, update the activity record and PR metadata, and leave only the final charge insertion/rebuild for settlement.

Assistant

2026-07-19 09:14:05.993 UTC

Role: assistant · Checkpoint: CP-013 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

The independent work is now complete: all three attempts have rendered learning curves, firing-rate panels, and matched rasters; the report builds with them; the artifact integrity/privacy checks pass; GitHub preview and tests pass. CI lint remains only the previously documented unrelated type debt. I'm checkpointing this progress—including your status exchange—while final billing remains pending.

Checkpoint CP-014

Provider reconciliation and final atomic publication are complete. This checkpoint follows diagnostics commit `3e27721`. Its immutable private source has SHA-256 prefix `c4b4c0583be4`, and its checkpoint time is `2026-07-19 09:17:58.200 UTC`.

Decision and action ledger

1. Candidate 3's authoritative charge is 0.08879619 USD: 0.04449077 USD COBA and 0.04430542 USD PING.
2. Exact cumulative exp070 spend is 0.47976851 USD across all six pods. The provider ledgers and aggregate `numbers.json` report zero active pods and confirm that every recorded attempt has exact billing.
3. Candidate 3's complete raw metrics, matched inputs, rasters, curves, firing diagnostics, decision, smoke, configuration, and ledger are atomically published. The final publication manifest is clean.
4. Pending work is limited to final validation, Demolab build, commit/push, preview verification, PR review-gate comment, and human handoff. No merge is authorized.

Visible messages added in CP-014

Assistant

2026-07-19 09:16:07.871 UTC

Role: assistant · Checkpoint: CP-014 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Good timing—the final rows have landed. Candidate 3 cost exactly \$0.04449077 COBA + \$0.04430542 PING = \$0.08879619. Exact cumulative exp070 spend is \$0.47976851, with zero active pods. I'm doing the final atomic publication and reconciliation now.

Assistant

2026-07-19 09:17:49.966 UTC

Role: assistant · Checkpoint: CP-014 · Session: 019f76ab-932b-72c2-b163-463ead65d8d5

Final reconciliation is complete. All three attempt ledgers, raw metrics, matched inputs, rasters, curves, activity diagnostics, and the aggregate `numbers.json` are now atomically published from a clean commit state. I'm adding the final hash-linked activity checkpoint, then running the complete validation/build/push/preview gate.