

Strong vs Weak Coupling

exp060 · 26 June 2026 · Draft

What does it mean to say a recurrent network is *strongly coupled*? It is tempting to read it as “dense wiring” or “big synaptic weights,” but neither alone is the definition. The real answer is about how the synaptic weight *scales* as the network grows — and getting it right is what separates a network that fires like cortex from one that does not. This note derives that from scratch, and the two figures are computed straight from the scaling laws, not simulated.

The short version

- **It is about scaling, not size.** Strong coupling means the per-synapse weight shrinks *slowly* — as $1/\sqrt{K}$ — as the in-degree K grows. Weak coupling shrinks it faster, as $1/K$.
- **Strong coupling forces balance.** The $1/\sqrt{K}$ scaling makes each population’s gross input $O(\sqrt{K})$ — enormous — so excitation and inhibition have no choice but to cancel. The small $O(1)$ residual that is left over, dominated by fluctuations, drives irregular ($CV \approx 1$), cortex-like firing.
- **Weak coupling kills the fluctuations.** The $1/K$ scaling sends the input fluctuation to zero; the network goes smooth and deterministic — analytically convenient, but not cortex.
- **Weight is not coupling.** The weight w you set and the coupling $J = w\sqrt{K}$ that fixes the regime differ by a factor of $\sqrt{K}$; “strongly coupled” is always a statement about J , never w alone.

The rest of the note earns each of these.

What one neuron receives

Take a single neuron in a large network. The input it integrates from one presynaptic population has a mean and a fluctuating part:

$$\mu \approx Kwr \quad (1)$$

$$\sigma \approx \sqrt{Kw^2r} = \sqrt{K}w\sqrt{r} \quad (2)$$

where

- K — the number of presynaptic inputs onto the neuron (its in-degree), set by connection density times population size; here $K = (1 - s)N_{\text{pre}}$, with s the sparsity.
- w — the **per-synapse weight**: the mean of the recurrent weight matrix W , i.e. one synapse’s peak conductance, the value you pass to the simulator ($-w\text{-ei}$ and friends).
- r — the mean firing rate of the presynaptic population.
- μ, σ — the mean and standard deviation of the total input from that population (the K spike trains treated as independent, so variances add).

Equation (1) is “ K inputs, each of size w , each at rate r .” Equation (2) is the same sum’s fluctuation. Everything that follows hinges on one question: **as the network grows, how does w scale with K ?** Two answers, two worlds.

Weak coupling — weights shrink as $1/K$

The classical mean-field choice keeps the mean input finite by shrinking each weight in step with the number of inputs:

$$w = \frac{w_0}{K} \implies \mu \approx w_0r = O(1), \quad \sigma \approx \frac{w_0\sqrt{r}}{\sqrt{K}} \xrightarrow{K \rightarrow \infty} 0 \quad (3)$$

Each input is a tiny nudge, the mean is whatever the drive makes it, and the fluctuations *vanish* as the network grows. The neuron rides the population average, smoothly and deterministically — the regime behind classical rate models and Wilson–Cowan dynamics. Convenient, but it cannot produce the irregular, Poisson-like firing of cortex.

Strong coupling — weights shrink as $1/\sqrt{K}$

The van Vreeswijk–Sompolinsky choice shrinks the weights more slowly:

$$w = \frac{J}{\sqrt{K}} \implies \mu \approx \sqrt{K} Jr = O(\sqrt{K}), \quad \sigma \approx J\sqrt{r} = O(1) \quad (4)$$

Now two things happen at once. The fluctuation stays finite, $O(1)$ — it does not wash out, and it is the seed of irregular firing. But the *mean* input from a single population is now $O(\sqrt{K})$, enormous for a realistic cell ($K \sim 10^3$ – 10^4). A current that large would pin the neuron far above threshold and saturate its rate. Something has to cancel it — which is the whole story, two sections down.

Weight is not coupling

That constant J in equation (4) is the **coupling**, and it is worth pinning down because it is constantly confused with the weight:

- the **per-synapse weight** w (the matrix mean W , what you set) is a number you choose;
- the **coupling** $J = w\sqrt{K}$ is *derived* — it tells you which regime that choice puts you in.

They differ by a factor of $\sqrt{K}$, and that is the point. Two networks with the same weight but different fan-in are coupled differently; two with the same J are in the same regime even if their weights and densities differ. So “is it strongly coupled?” is a question about J (equivalently, about how w compares to $1/\sqrt{K}$), never about w alone. Holding J fixed as $K \rightarrow \infty$ *means* letting the weight shrink as $w = J/\sqrt{K}$ — that is the operating definition of strong coupling.

(Cross-referencing note: exp058 currently writes this relation inverted, $J = w/\sqrt{K}$. The version here — $J = w\sqrt{K}$, weights shrinking as $1/\sqrt{K}$ — is the one that keeps σ at $O(1)$, so it is the convention to trust; the two are being reconciled.)

pending re-run with new canonical data

Figure 1: Equations (3)–(4) made visual. **Left:** the mean recurrent input μ — under strong coupling it climbs $\propto \sqrt{K}$ (so each population’s drive is huge and must cancel); under weak coupling it is flat. **Right:** the fluctuation σ , and the punchline — under strong coupling it stays $O(1)$ as the network grows (irregular firing survives), while under weak coupling it decays $\propto 1/\sqrt{K}$ to zero (the network goes deterministic). Analytic; no simulation.

Why strong coupling forces balance

The only thing that can cancel an $O(\sqrt{K})$ excitatory current is an equally large $O(\sqrt{K})$ inhibitory one. With both populations:

$$\mu_{\text{net}} \approx \sqrt{K}(J_E r_E - J_I r_I) \quad (5)$$

For the neuron to fire at a finite rate, the bracket must be small — of order $1/\sqrt{K}$ — so that μ_{net} returns to $O(1)$:

$$J_E r_E - J_I r_I = O\left(\frac{1}{\sqrt{K}}\right) \quad (6)$$

This is the **balance condition**, and the key point is that *nobody imposes it by hand*. You cannot tune external parameters to satisfy (6) at every K ; instead the rates r_E and r_I adjust themselves until gross excitation and gross inhibition track and nearly cancel. The cancellation is self-consistent and emergent. What is left over — the small residual between two huge opposing currents — is dominated by the $O(1)$ fluctuations of (4), not a steady drive. So the membrane potential hovers near threshold and crosses it at irregular, fluctuation-driven moments. (r_E, r_I are the self-organised population rates; J_E, J_I the two couplings; μ_{net} the net input, a small difference of two $O(\sqrt{K})$ terms.)

What the balanced state gives you

Because firing is driven by the residual rather than a mean overshoot, three cortically realistic things follow:

- **Irregular spiking.** Inter-spike intervals look Poisson, $CV \approx 1$ — the hallmark of cortical firing — with no injected noise. The irregularity is intrinsic and, in the deterministic limit, chaotic.
- **Fast, linear tracking.** Balance is dynamic (rates re-cancel almost instantly), so the network responds on a single-synapse timescale and the population rate is roughly linear in the drive.
- **Robustness, not fine-tuning.** Balance emerges from the strong-coupling scaling itself; it is not a knife-edge.

This is why “strongly coupled, recurrent, balanced” travel together: strong coupling ($w \sim 1/\sqrt{K}$) *forces* balance, and balance makes the state irregular and fast. The simulated counterpart of Figure 1 — irregularity (CV) surviving as K grows under strong coupling but decaying under weak — is measured directly in exp058, Figure 2.

The conductance-based wrinkle

Equations (1)–(6) treat inputs as currents. Real synapses — and this project’s model — are conductance-based: a synapse opens a conductance g and injects $g(V - E_{\text{rev}})$, which depends on the voltage and reverses at E_{rev} . Strong coupling then means the *total* synaptic conductance is large, with a consequence the current picture misses — the effective membrane time constant shrinks,

$$\tau_{\text{eff}} = \frac{C_m}{g_L + g_{\text{tot}}} \quad (7)$$

so a strongly coupled conductance-based neuron is not just driven harder, it integrates *faster* and its gain changes with the input level. That is why conductance-based balanced networks need their own theory: the convenient current-based scaling laws mispredict rates and correlations once g_{tot} ($O(\sqrt{K})$ here) is large. (C_m membrane capacitance, g_L leak conductance, E_{rev} reversal potential.)

Putting it in this project's terms

The model used throughout this lab lives in the strong-coupling world by construction, and four quantities chain together. Sparsity sets the fan-in; fan-in and weight set the coupling:

$$K = (1 - s)N_{\text{pre}}, \quad J = w\sqrt{K} = w\sqrt{(1 - s)N_{\text{pre}}} \quad (8)$$

You choose s and N_{pre} (which fix K) and you choose w ; together w and K land you in regime J . So the two knobs you actually turn — s and w — are *not* independent. Making the network sparser at fixed w *weakens* the coupling, and to stay in the same balanced regime you must raise the weight to compensate:

$$w \propto \frac{1}{\sqrt{(1 - s)N_{\text{pre}}}} \quad (\text{to hold } J \text{ fixed}) \quad (9)$$

pending re-run with new canonical data

Figure 2: Equation (9). To stay in the same coupling regime as the network is made sparser (right, smaller K), the per-synapse weight has to grow — gently at first, then steeply as $s \rightarrow 1$. Sparsity s and weight w are not independent: a change in one must be matched in the other to hold J fixed.

So W is the weight you set, s sets the fan-in K , and J is the regime the pair lands you in. The balanced state itself — irregular, asynchronous, rhythmless firing with $CV \approx 1$ — is built and measured in exp058, where *weights scaled to hold J fixed* versus *weights held fixed* is exactly the strong-versus-weak distinction made here. PING then sits one step beyond: it keeps the strong, balanced E–I coupling but adds a *structured* recurrent loop (E→I→E) that can turn the asynchronous balanced regime into a gamma rhythm — the onset studied in the manuscript. Strong coupling is the substrate; the structured loop is what makes it tick.

Sources

- van Vreeswijk & Sompolinsky (1996), *Chaos in Neuronal Networks with Balanced Excitatory and Inhibitory Activity* — the origin of the $1/\sqrt{K}$ strong-coupling scaling and the balanced state.
- Brunel (2000), *Dynamics of Sparsely Connected Networks of Excitatory and Inhibitory Spiking Neurons* — the analytic phase diagram of where these regimes live and oscillate.
- Renart et al. (2010), *The Asynchronous State in Cortical Circuits* — how balance also decorrelates the population, not just individual cells.
- Gerstner (2000), *Population Dynamics of Spiking Neurons* — the population-level companion.