

Switching the loop on at inference cuts E rate $\approx 10\times$

nb038 · 30 May 2026

The trained networks this entry uses are produced once in the shared training hub, nb022 (Training), and reused here rather than retrained.

Abstract

Loads a trained COBA network and switches the I-loop on at inference by sweeping *ei_strength* from 0 to 1. The same feedforward weights that fire at ≈ 52 Hz without the loop fire at ≈ 5 Hz with the loop engaged, and accuracy stays within ≈ 12 pp of the COBA baseline. PING gating is a post-hoc sparsity knob — the architecture, not the training, supplies the gamma dynamics.

Method

Parameter	Value
Integration timestep Δt	0.1 ms
Trial duration T	200 ms
MNIST samples (80/20 stratified split of 2000)	1600 train / 400 test ($\approx 2.9\%$ of the 70k-sample MNIST corpus)
Epochs	10

The PING and COBA baseline definitions and the training recipe are in nb025 this entry runs the COBA $\rightarrow$ PING I-loop transfer probe at eval time on the trained baselines. (The input-rate sweep / f-I curve material that previously lived here has moved to nb023, the natural home for “architectural response to drive”).

Inference-time probe. The trained COBA baseline (seed 42, $\theta_u = \text{off}$) is loaded and *ei_strength* — the I-loop gain — is overridden at eval time across 11 values from 0 to 1. W_{in} and W_{out} load from the COBA checkpoint; W^{EI} and W^{IE} are freshly initialised (the COBA checkpoint stores these at zero, so skipping the load leaves a functional I-loop). No retraining.

Results




Figure 109: Switching the recurrent I-loop on *at inference* on a trained COBA network — no retraining, W^{EI}/W^{IE} freshly wired (the COBA checkpoint stores them at zero). **Top** — the same feedforward weights fire densely and asynchronously at $ei = 0$ (COBA) and in gamma bands at $ei = 1$ (PING); the gamma dynamics come from the inhibitory architecture, not from training. **Bottom left** — E rate falls $\approx 10\times$ ($\approx 60 \rightarrow 5$ Hz) as the loop engages while I rises to ≈ 28 Hz; the suppression is continuous in loop strength. **Bottom right** — accuracy *degrades* without retraining, from the $\approx 87\%$ COBA baseline to $\approx 63\%$ at full strength (a ≈ 24 pp cost). So the architecture supplies the rate-gating for free, but using it well needs training *with* the loop (nb025): the sparsity is architectural, the accuracy is learned.




Figure 110: The full transition: trained COBA replayed at six inference-time $ei_strength$ values (same trial, same feedforward weights, a fresh I-loop each row). At $ei = 0$ the asynchronous-dense COBA pattern persists; by $ei \approx 0.4$ the same weights produce gamma cycles, sharpening toward $ei = 1$. The rhythm appears continuously as the loop is wired in — no retraining at any point.