# Strong vs Weak Coupling

nb060 · 26 June 2026 · final · Miscellaneous

What does it mean to say a recurrent network is *strongly coupled*? It is tempting to read it as "dense wiring" or "big synaptic weights," but neither alone is the definition. The real answer is about how the synaptic weight *scales* as the network grows — and getting it right is what separates a network that fires like cortex from one that does not. This note derives that from scratch, and the two figures are computed straight from the scaling laws, not simulated.

## The short version

- **It is about scaling, not size.** Strong coupling means the per-synapse weight shrinks *slowly* — as $1/\sqrt{K}$ — as the in-degree $K$ grows. Weak coupling shrinks it faster, as $1/K$.
- **Strong coupling forces balance.** The $1/\sqrt{K}$ scaling makes each population's gross input $O\left(\sqrt{K}\right)$ — enormous — so excitation and inhibition have no choice but to cancel. The small $O(1)$ residual that is left over, dominated by fluctuations, drives irregular ($CV \approx 1$), cortex-like firing.
- **Weak coupling kills the fluctuations.** The $1/K$ scaling sends the input fluctuation to zero; the network goes smooth and deterministic — analytically convenient, but not cortex.
- **Weight is not coupling.** The weight $w$ you set and the coupling $J = w\sqrt{K}$ that fixes the regime differ by a factor of $\sqrt{K}$; "strongly coupled" is always a statement about $J$, never $w$ alone.

The rest of the note earns each of these.

## What one neuron receives

Take a single neuron in a large network. The input it integrates from one presynaptic population has a mean and a fluctuating part:

$$\mu \approx Kwr \quad (1)$$

$$\sigma \approx \sqrt{Kw^2r} = \sqrt{K}w\sqrt{r} \quad (2)$$

where

- $K$ — the number of presynaptic inputs onto the neuron (its in-degree), set by connection density times population size; here $K = (1 - s)N_{\text{pre}}$, with $s$ the sparsity.
- $w$ — the **per-synapse weight**: the mean of the recurrent weight matrix $W$, i.e. one synapse's peak conductance, the value you pass to the simulator (*–w-ei* and friends).
- $r$ — the mean firing rate of the presynaptic population.
- $\mu$, $\sigma$ — the mean and standard deviation of the total input from that population (the $K$ spike trains treated as independent, so variances add).

Equation (1) is "K inputs, each of size $w$, each at rate $r$." Equation (2) is the same sum's fluctuation. Everything that follows hinges on one question: **as the network grows, how does $w$ scale with $K$?** Two answers, two worlds.

## Weak coupling — weights shrink as $1/K$

The classical mean-field choice keeps the mean input finite by shrinking each weight in step with the number of inputs:

$$w = \frac{w_0}{K} \quad \Longrightarrow \quad \mu \approx w_0 r = O(1), \quad \sigma \approx \frac{w_0\sqrt{r}}{\sqrt{K}} \underset{K\to\infty}{\to} 0 \quad (3)$$

Each input is a tiny nudge, the mean is whatever the drive makes it, and the fluctuations *vanish* as the network grows. The neuron rides the population average, smoothly and deterministically — the regime behind classical rate models and Wilson–Cowan dynamics. Convenient, but it cannot produce the irregular, Poisson-like firing of cortex.

## Strong coupling — weights shrink as $1/\sqrt{K}$

The van Vreeswijk–Sompolinsky choice shrinks the weights more slowly:

$$w = \frac{J}{\sqrt{K}} \quad \Longrightarrow \quad \mu \approx \sqrt{K}Jr = O\left(\sqrt{K}\right), \quad \sigma \approx J\sqrt{r} = O(1) \quad (4)$$

Now two things happen at once. The fluctuation stays finite, $O(1)$ — it does not wash out, and it is the seed of irregular firing. But the *mean* input from a single population is now $O\left(\sqrt{K}\right)$, enormous for a realistic cell ($K \sim 10^3$–$10^4$). A current that large would pin the neuron far above threshold and saturate its rate. Something has to cancel it — which is the whole story, two sections down.

## Weight is not coupling

That constant $J$ in equation (4) is the **coupling**, and it is worth pinning down because it is constantly confused with the weight:

- the **per-synapse weight** $w$ (the matrix mean $W$, what you set) is a number you choose;
- the **coupling** $J = w\sqrt{K}$ is *derived* — it tells you which regime that choice puts you in.

They differ by a factor of $\sqrt{K}$, and that is the point. Two networks with the same weight but different fan-in are coupled differently; two with the same $J$ are in the same regime even if their weights and densities differ. So “is it strongly coupled?” is a question about $J$ (equivalently, about how $w$ compares to $1/\sqrt{K}$), never about $w$ alone. Holding $J$ fixed as $K \to \infty$ *means* letting the weight shrink as $w = J/\sqrt{K}$ — that is the operating definition of strong coupling.

(Cross-referencing note: nb058 currently writes this relation inverted, $J = w/\sqrt{K}$. The version here — $J = w\sqrt{K}$, weights shrinking as $1/\sqrt{K}$ — is the one that keeps $\sigma$ at $O(1)$, so it is the convention to trust; the two are being reconciled.)

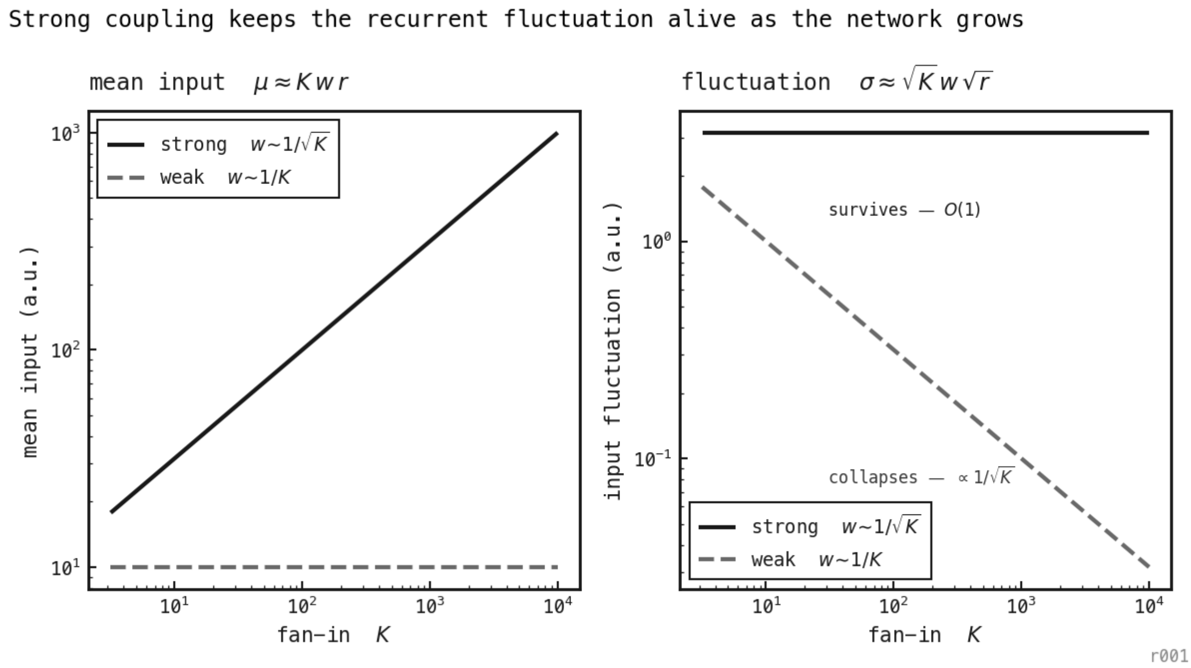


Figure 177: Equations (3)–(4) made visual. **Left:** the mean recurrent input $\mu$ — under strong coupling it climbs $\propto \sqrt{K}$ (so each population's drive is huge and must cancel); under weak coupling it is flat. **Right:** the fluctuation $\sigma$, and the punchline — under strong coupling it stays $O(1)$ as the network grows (irregular firing survives), while under weak coupling it decays $\propto 1/\sqrt{K}$ to zero (the network goes deterministic). Analytic; no simulation.

## Why strong coupling forces balance

The only thing that can cancel an $O(\sqrt{K})$ excitatory current is an equally large $O(\sqrt{K})$ inhibitory one. With both populations:

$$\mu_{\text{net}} \approx \sqrt{K}(J_E r_E - J_I r_I) \quad (5)$$

For the neuron to fire at a finite rate, the bracket must be small — of order $1/\sqrt{K}$ — so that $\mu_{\text{net}}$ returns to $O(1)$:

$$J_E r_E - J_I r_I = O\left(\frac{1}{\sqrt{K}}\right) \quad (6)$$

This is the **balance condition**, and the key point is that *nobody imposes it by hand.* You cannot tune external parameters to satisfy (6) at every $K$; instead the rates $r_E$ and $r_I$ adjust themselves until gross excitation and gross inhibition track and nearly cancel. The cancellation is self-consistent and emergent. What is left over — the small residual between two huge opposing currents — is dominated by the $O(1)$ fluctuations of (4), not a steady drive. So the membrane potential hovers near threshold and crosses it at irregular, fluctuation-driven moments. ($r_E, r_I$ are the self-organised population rates; $J_E, J_I$ the two couplings; $\mu_{\text{net}}$ the net input, a small difference of two $O(\sqrt{K})$ terms.)

## What the balanced state gives you

Because firing is driven by the residual rather than a mean overshoot, three cortically realistic things follow:

- **Irregular spiking.** Inter-spike intervals look Poisson, $CV \approx 1$ — the hallmark of cortical firing — with no injected noise. The irregularity is intrinsic and, in the deterministic limit, chaotic.
- **Fast, linear tracking.** Balance is dynamic (rates re-cancel almost instantly), so the network responds on a single-synapse timescale and the population rate is roughly linear in the drive.
- **Robustness, not fine-tuning.** Balance emerges from the strong-coupling scaling itself; it is not a knife-edge.

This is why "strongly coupled, recurrent, balanced" travel together: strong coupling ($w \sim 1/\sqrt{K}$) *forces* balance, and balance makes the state irregular and fast. The simulated counterpart of Figure 1 — irregularity ($CV$) surviving as $K$ grows under strong coupling but decaying under weak — is measured directly in nb058, Figure 2.

## The conductance-based wrinkle

Equations (1)–(6) treat inputs as currents. Real synapses — and this project's model — are conductance-based: a synapse opens a conductance $g$ and injects $g(V - E_{\text{rev}})$, which depends on the voltage and reverses at $E_{\text{rev}}$. Strong coupling then means the *total* synaptic conductance is large, with a consequence the current picture misses — the effective membrane time constant shrinks,

$$\tau_{\text{eff}} = \frac{C_m}{g_L + g_{\text{tot}}} \quad (7)$$

so a strongly coupled conductance-based neuron is not just driven harder, it integrates *faster* and its gain changes with the input level. That is why conductance-based balanced networks need their own theory: the convenient current-based scaling laws mispredict rates and correlations once $g_{\text{tot}}$ ($O\left(\sqrt{K}\right)$ here) is large. ($C_m$ membrane capacitance, $g_L$ leak conductance, $E_{\text{rev}}$ reversal potential.)

## Putting it in this project's terms

The model used throughout this lab lives in the strong-coupling world by construction, and four quantities chain together. Sparsity sets the fan-in; fan-in and weight set the coupling:

$$K = (1 - s)N_{\text{pre}}, \quad J = w\sqrt{K} = w\sqrt{(1 - s)N_{\text{pre}}} \quad (8)$$

You choose $s$ and $N_{\text{pre}}$ (which fix $K$) and you choose $w$; together $w$ and $K$ land you in regime $J$. So the two knobs you actually turn — $s$ and $w$ — are *not* independent. Making the network sparser at fixed $w$ *weakens* the coupling, and to stay in the same balanced regime you must raise the weight to compensate:

$$w \propto \frac{1}{\sqrt{(1 - s)N_{\text{pre}}}} \quad \text{(to hold } J \text{ fixed)} \quad (9)$$

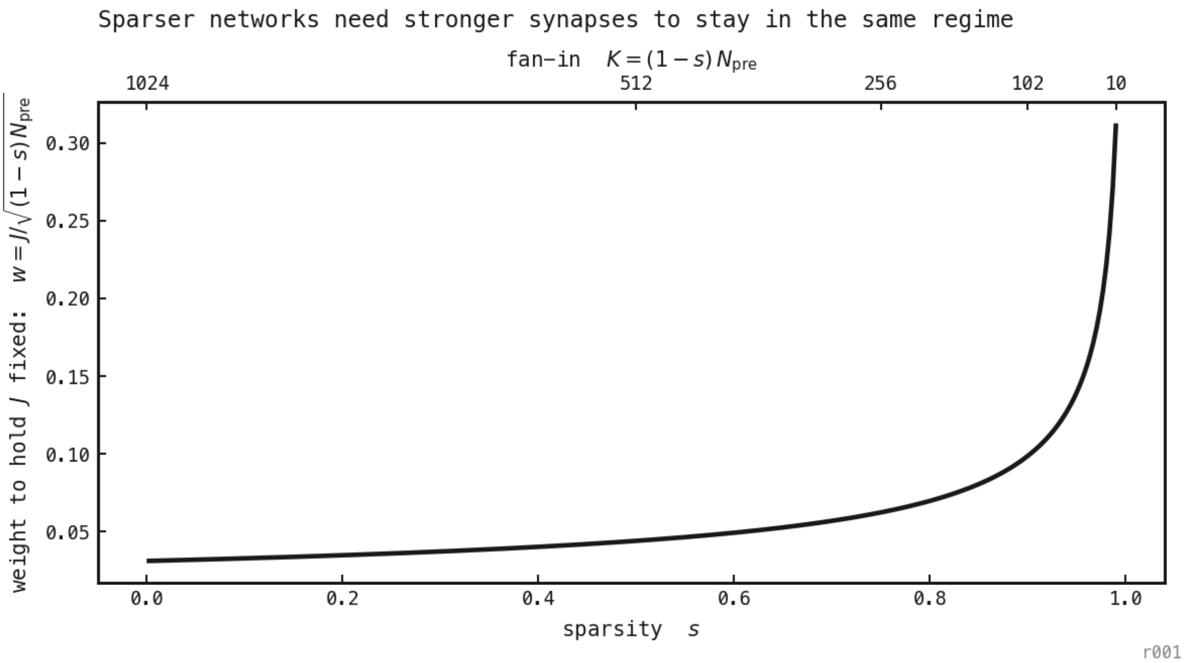


Figure 178: Equation (9). To stay in the same coupling regime as the network is made sparser (right, smaller $K$), the per-synapse weight has to grow — gently at first, then steeply as $s \to 1$. Sparsity $s$ and weight $w$ are not independent: a change in one must be matched in the other to hold $J$ fixed.

So $W$ is the weight you set, $s$ sets the fan-in $K$, and $J$ is the regime the pair lands you in. The balanced state itself — irregular, asynchronous, rhythmless firing with $CV \approx 1$ — is built and measured in nb058, where *weights scaled to hold J fixed* versus *weights held fixed* is exactly the strong-versus-weak distinction made here. PING then sits one step beyond: it keeps the strong, balanced E–I coupling but adds a *structured* recurrent loop (E→I→E) that can turn the asynchronous balanced regime into a gamma rhythm — the onset studied in the manuscript. Strong coupling is the substrate; the structured loop is what makes it tick.

## Sources


- van Vreeswijk & Sompolinsky (1996), *Chaos in Neuronal Networks with Balanced Excitatory and Inhibitory Activity* — the origin of the $1/\sqrt{K}$ strong-coupling scaling and the balanced state.
- Brunel (2000), *Dynamics of Sparsely Connected Networks of Excitatory and Inhibitory Spiking Neurons* — the analytic phase diagram of where these regimes live and oscillate.
- Renart et al. (2010), *The Asynchronous State in Cortical Circuits* — how balance also decorrelates the population, not just individual cells.
- Gerstner (2000), *Population Dynamics of Spiking Neurons* — the population-level companion, reviewed in ar059.